IsoColBERT: Token-Level Isotropy Regularization for Multilingual Late-Interaction Retrieval

Isaac Chang Ingrid Chien M. Tran | Harvard University / MIT | MRL Workshop @ EMNLP 2026

Late-interaction multilingual retrievers such as ColBERT use token-level embeddings for MaxSim matching. We show that multilingual ColBERT models **suffer from token-level anisotropy**, where embeddings collapse into a narrow cone with low effective rank. To address this, we introduce **IsoColBERT**, a regularizer added to the InfoNCE training objective with zero inference cost. Across three random seeds, IsoColBERT improves token-embedding geometry and significantly boosts low-resource retrieval, achieving a **+2.69 point** gain on English–Swahili ($p < 0.001$). Ablations show that broad embedding-space regularization is more effective than localized token- and attention-based constraints. These results demonstrate that token-level isotropy directly impacts multilingual late-interaction retrieval.

Motivation

Cross-lingual search and retrieval-augmented generation depend on multilingual retrievers — systems that map English queries to documents in dozens of target languages. Late-interaction models like **ColBERT** differ from sentence-level encoders by matching at the **token level**: every query token finds its best partner in the candidate document, and the scores are summed (the *MaxSim* operation). Retrieval quality therefore depends directly on the **geometry of individual token vectors**.

Why this matters now.

- In *single-vector* encoders, anisotropy is averaged out across the sequence before scoring — its impact is muted.
- In *late-interaction* models, **every token participates directly**. If embeddings collapse into a narrow cone, MaxSim degenerates: many tokens look similar to many other tokens.

Anisotropy in single-vector encoders is well documented (Ethayarajh, EMNLP 2019). Its consequences for multi-vector retrievers remain largely unexamined — yet token geometry is precisely what these models rely on at inference time.

Intuition. Imagine packing arrows into a tiny cone vs. spreading them across a sphere. Both have the same number of arrows, but the sphere lets each one "point uniquely" — which is exactly what MaxSim needs to distinguish good matches from bad ones.

The pathology should hit hardest on **low-resource cross-lingual pairs**, where the encoder has the least pretrained signal and the largest retrieval headroom to recover.

Method: IsoColBERT

Standard ColBERT setup: XLM-R encoder $\rightarrow$ linear projection to 128 dims $\rightarrow$ L_2-normalized token vectors for queries and documents.

MaxSim retrieval score (sums per-query-token best matches):

$$\text{score}(Q, D) = \sum_{q \in Q} \max_{d \in D} \langle q, d \rangle$$

Training loss (IsoColBERT):

$$\mathcal{L} = \mathcal{L}_{\text{InfoNCE}} + \lambda \cdot \frac{1}{2}\left(\overline{\cos^2}_{\text{off}}(Q) + \overline{\cos^2}_{\text{off}}(D)\right)$$

The second term is the **mean squared off-diagonal cosine similarity** among valid query and document token vectors in a batch. It penalizes any pair of tokens that look too similar, regardless of their content.

What the loss does. Push every pair of token vectors slightly apart. Over training, the model learns to "use more of the sphere" — each token finds its own direction. Inference is unchanged: the regularizer only shapes training-time geometry.

One hyperparameter (λ, regularizer strength). **Zero inference cost.** No architecture changes.

Experimental Setup

Model. XLM-RoBERTa-base fine-tuned end-to-end with a 128-dim linear projection producing L_2-normalized token vectors. Sequences are truncated to 64 tokens. We train with the combined InfoNCE and isotropy objective; the baseline uses InfoNCE alone at $\lambda=0$.
Training. OPUS-100 English–Spanish parallel text. We train for 2000 steps at effective batch size 32 using AdamW with learning rate 2×10^{-5}, weight decay 0.01, and InfoNCE temperature 0.07.
Evaluation. FLORES+ dev and devtest are combined into a single per-language retrieval pool of roughly 2009 candidates. We evaluate across five language pairs: EN$\rightarrow$ES, EN$\rightarrow$FR, and EN$\rightarrow$DE as high-resource controls; EN$\rightarrow$SW for Swahili and EN$\rightarrow$AR for Arabic as low-resource headline pairs. All pairs except EN$\rightarrow$ES are evaluated **zero-shot**: the model never sees these target languages during contrastive training.
Metrics.

- *MaxSim Precision@1.* Primary retrieval metric; full pairwise MaxSim scoring against the entire candidate pool.
- *Token intra-cosine.* Mean off-diagonal cosine across 2048 valid token vectors per language; direct measure of anisotropy.
- *Effective rank.* Exp of entropy of the normalized singular-value spectrum over 4096 stacked token vectors; how many dimensions the embedding space effectively uses.

Robustness. All retrieval results are reported as mean $\pm$ standard deviation across 3 independent random seeds. The λ ablation was conducted as a single-seed pilot; the best λ was then validated multi-seed.

Main Results: Retrieval

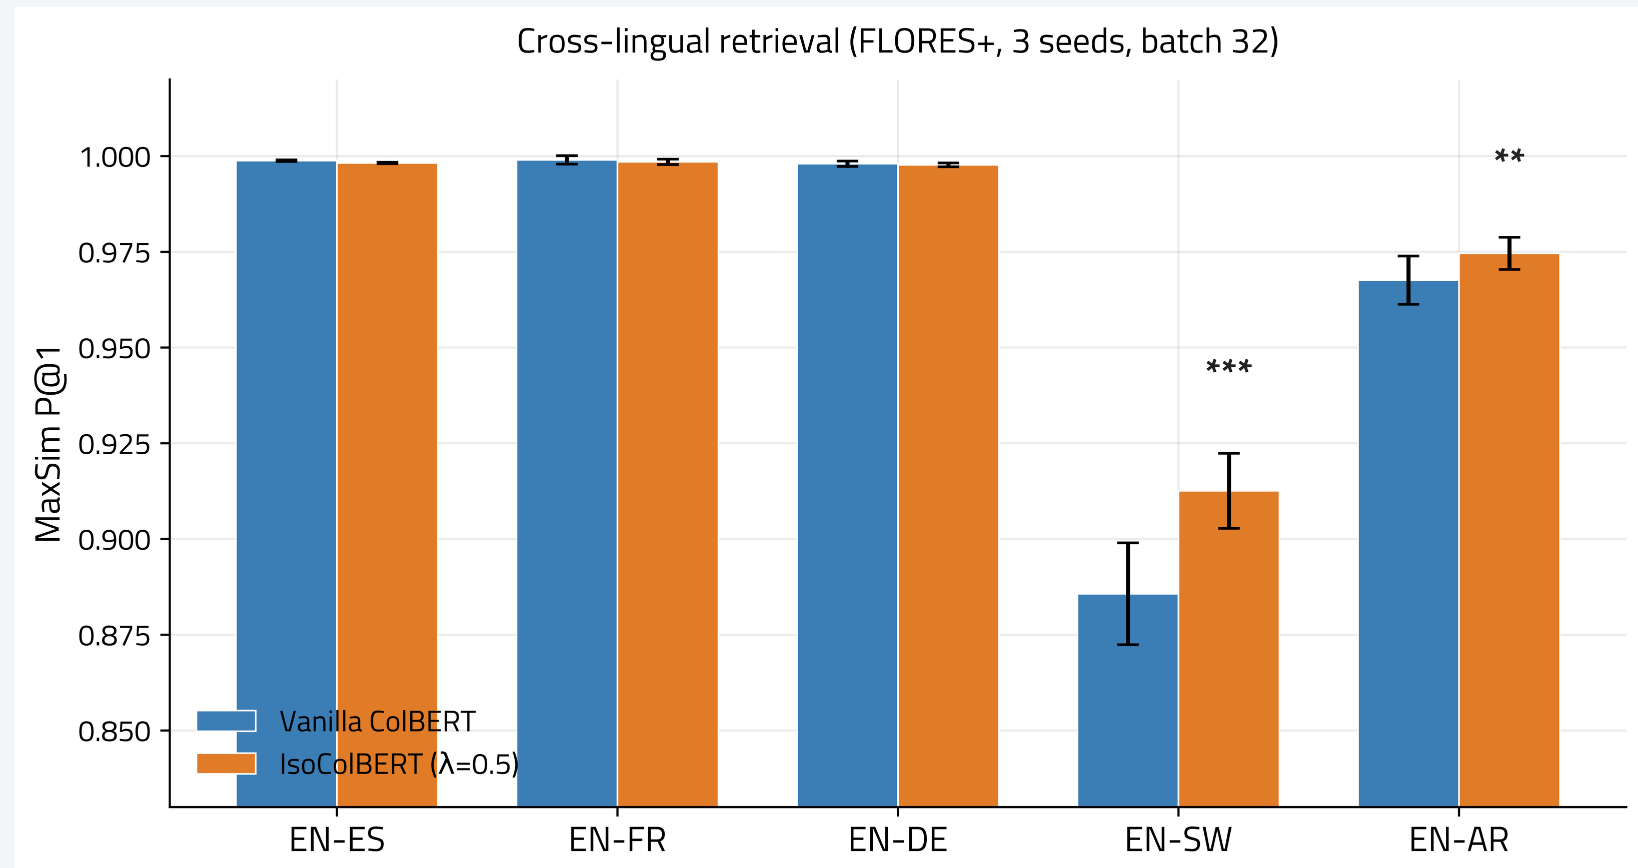

Headline. +2.69 pp on EN–SW ($p < 0.001$, all 3 seeds individually positive). **+0.70 pp** on EN–AR ($p < 0.01$). High-resource pairs (ES, FR, DE) saturate at $\geq$0.997 P@1 for both methods — both already at ceiling, no retrieval headroom for either method to exploit.

Why these languages? EN–SW and EN–AR are the only pairs with room to grow. XLM-R was pretrained on far less Swahili / Arabic data than Spanish / French / German — which is exactly where token geometry has the most room to be reshaped, and the most retrieval headroom to recover.

Per-seed sanity (EN–SW): seed 42 +0.018, seed 1337 +0.023, seed 2024 +0.039. Every seed individually positive, and variance also shrinks ($\sigma : 0.013 \rightarrow 0.010$). The regularized model is not just better on average — it is more stable.

Main Results: Geometry

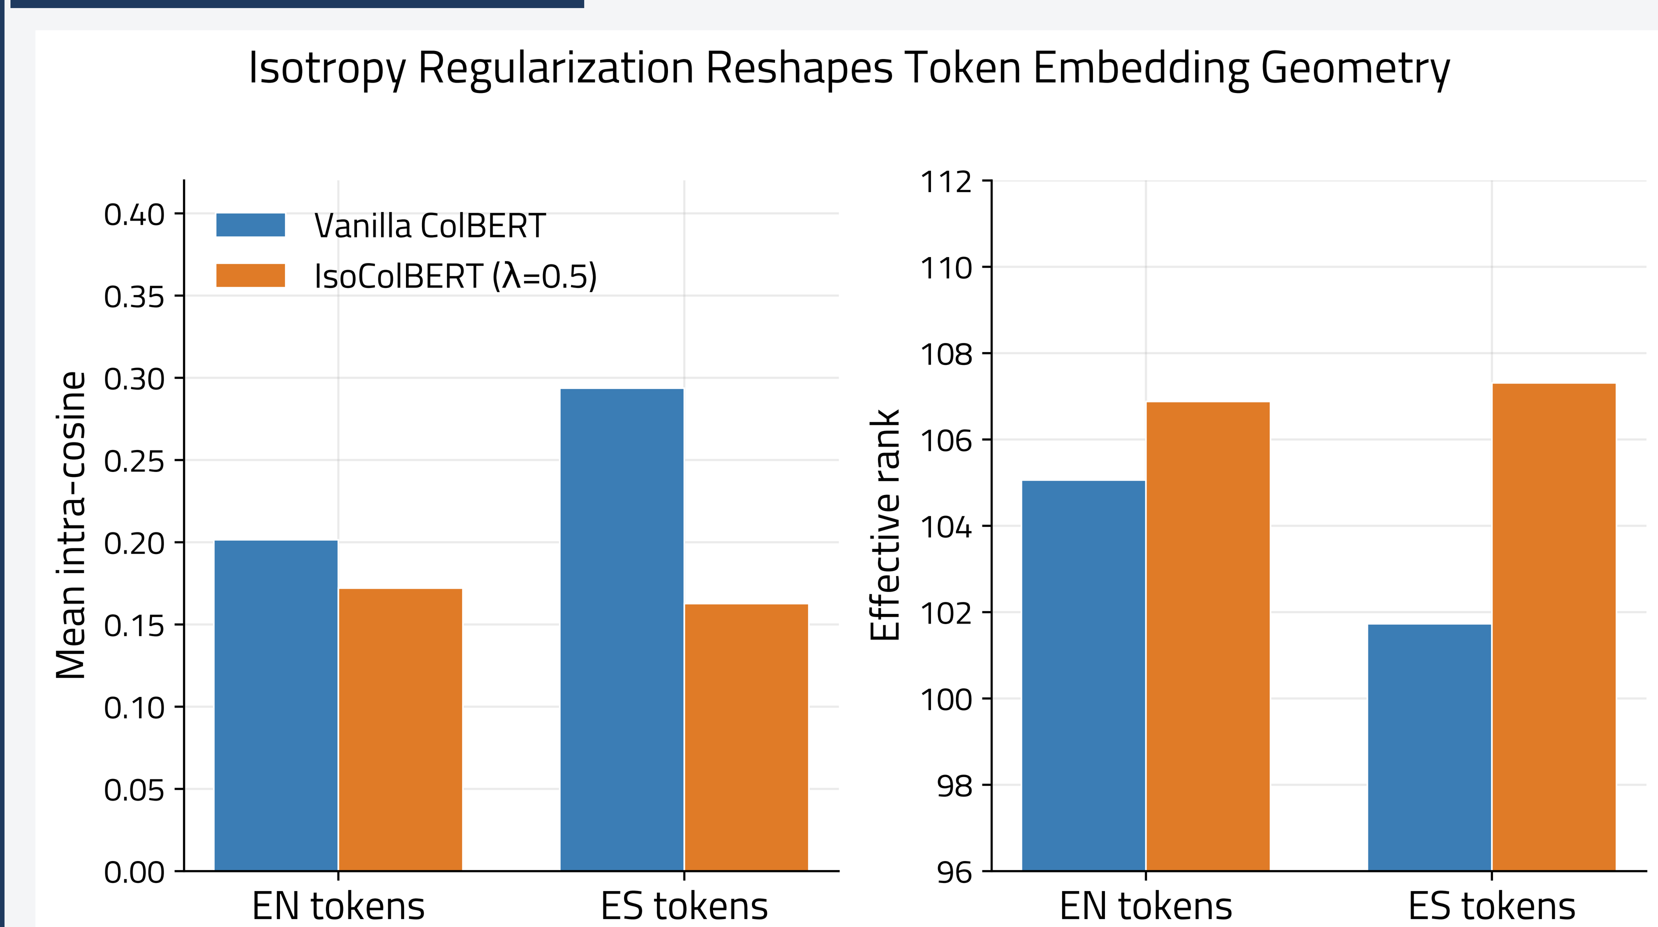

The regularizer reshapes the embedding space in exactly the direction it was designed to: **tokens become less similar to each other** (lower intra-cosine), and the embedding space **uses more of its available dimensions** (higher effective rank). All four shifts are significant at $p < 0.001$.

Why is ES affected more than EN? The model trains on EN–ES bitext using bidirectional InfoNCE. The Spanish (target-side) representations bear most of the contrastive pressure, so they collapse the hardest under vanilla training — and reshape the hardest under IsoColBERT.

Hard-negative robustness. We re-evaluate against K=20 LaBSE-mined distractors per query (sentences that another strong multilingual model judges most similar to the query). The EN–SW gain *still* holds at **+2.16 pp** ($p < 0.001$) — the improvement is not an artifact of weak FLORES candidate pools.
Five metrics improve at $p < 0.001$ simultaneously: EN–SW P@1, intra-cos EN, intra-cos ES, eRank EN, eRank ES.

Mechanism Ablations

We test two alternative places to put the regularizer.

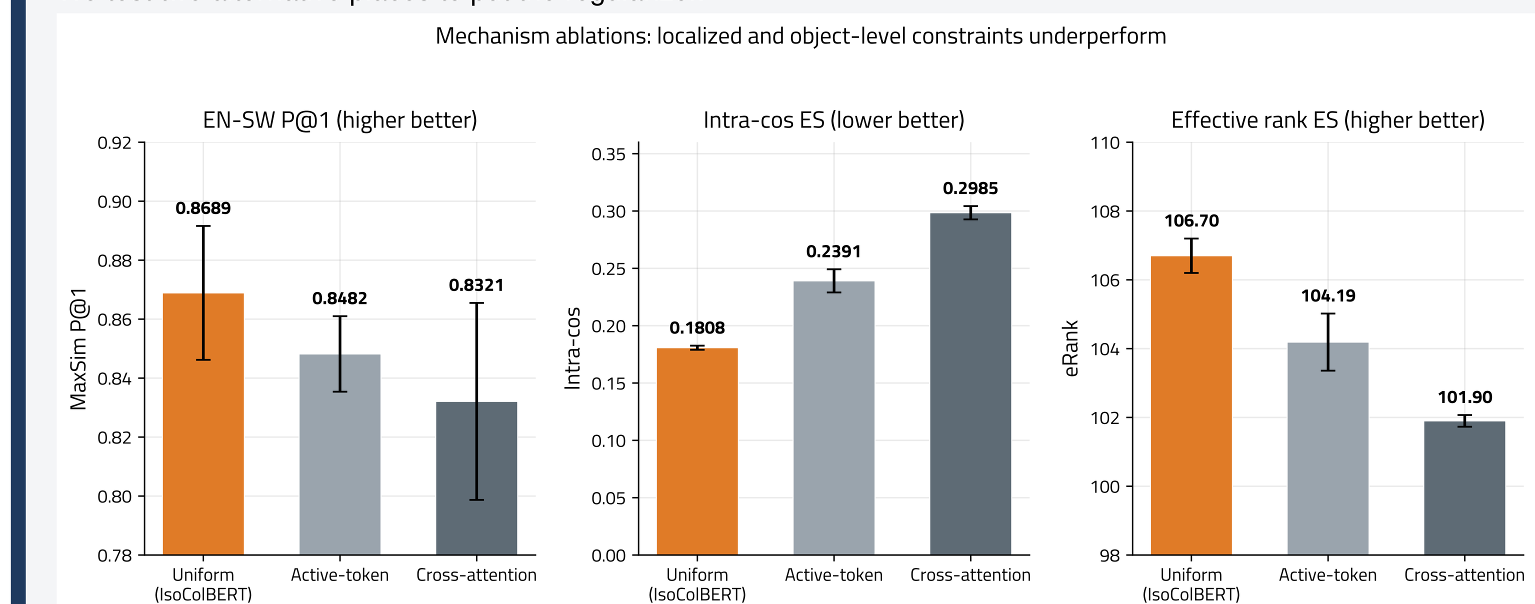

(a) Active-token iso-reg. Regularize *only* the tokens that win MaxSim matches (query tokens + their argmax doc tokens). Result: EN–SW -2.1 pp, intra-cos ES $+0.06$, eRank ES -2.5 vs. uniform.
(b) Cross-attention iso-reg. Regularize the $L_q \times L_d$ matching matrix itself — penalize redundancy in *matching patterns* rather than embedding directions. Result: EN–SW -3.7 pp, intra-cos ES $+0.12$, eRank ES -4.8 vs. uniform.

Why do these fail?
Active-token: Only fixes geometry for tokens that already win — but the retrieval loss comes from tokens that *lose* to a wrong-document token. Shrinking the regularization surface area kills the effect.
Cross-attention: Flattens the matching pattern but leaves the underlying tokens just as anisotropic as vanilla. Treats a symptom, not the cause.

Joint interpretation. Both negative results triangulate the conclusion: the benefit of iso-reg in late-interaction retrievers requires **broad embedding-space shaping** across the full token manifold. Localized or object-level constraints are not substitutes.

Lambda Ablation

Metric	Vanilla	$\lambda=0.1$	$\lambda=0.5$
EN–SW P@1	0.8344	0.8669	**0.8689**
EN–AR P@1	0.9423	0.9574	**0.9622**
Intra-cos ES	0.2937	0.2535	**0.1808**
eRank ES	101.73	103.78	**106.70**

Geometric metrics improve **monotonically** with λ. Retrieval P@1 begins saturating near $\lambda = 0.5$ — more pressure keeps reshaping the space, but the retrieval payoff plateaus.

Reading. $\lambda = 0$ is the vanilla collapsed state. Each step up in λ "spreads the sphere" further. Past $\lambda \approx 0.5$, you keep getting better geometry but no extra retrieval gain — the model has already escaped the regime where anisotropy was costing you accuracy.

Takeaways

- Token-level anisotropy is a **real, measurable pathology** in multilingual late-interaction retrievers — not just a single-vector problem.
- A **single regularizer term** fixes it with **zero inference cost** and no architecture changes.
- Gains concentrate where there is **both pretrained-signal scarcity and retrieval headroom** (low-resource pairs); high-resource pairs are unaffected — nothing breaks.
- Two negative ablations isolate the mechanism: **broad embedding-space shaping is the right level of intervention**.
- **Token geometry is a tractable lever** for low-resource multilingual retrieval.

Future work. Multi-pair training (mMARCO, MIRACL); other multilingual backbones; harder retrieval benchmarks beyond FLORES.